

Micro-Benchmarking MPI Partitioned Point-to-Point Communication

Yiltan Hassan Temuçin¹, Ryan E. Grant², and Ahmad Afsahi¹

¹The Parallel Processing Research Laboratory (PPRL)
Department of Electrical and Computer Engineering
Queen's University, Canada

²Computing at Extreme Scale Advanced Research Laboratory (CAESAR)
Department of Electrical and Computer Engineering
Queen's University, Canada

51st International Conference on Parallel Processing (ICPP '22),
August 29-September 1, 2022

Introduction

- Multi-Threaded MPI Point-to-Point Communication
- MPI Partitioned Point-to-Point Communication

Motivation

Micro-Benchmark Design And Results

- Overhead
- Perceived Bandwidth
- Application Availability
- Early Bird Communication
- Sweep3D Communication Pattern
- Halo3D Communication Pattern
- Potential Application Improvements

Conclusion And Future Work

- ▶ HPC is used to solve large complex problems in many domains
 - ▶ Communication is one of the main bottlenecks in applications

- ▶ HPC is used to solve large complex problems in many domains
 - ▶ Communication is one of the main bottlenecks in applications
- ▶ The Message Passing Interface (MPI)

- ▶ HPC is used to solve large complex problems in many domains
 - ▶ Communication is one of the main bottlenecks in applications
- ▶ The Message Passing Interface (MPI)
 - ▶ Popular parallel programming model in HPC

- ▶ HPC is used to solve large complex problems in many domains
 - ▶ Communication is one of the main bottlenecks in applications
- ▶ The Message Passing Interface (MPI)
 - ▶ Popular parallel programming model in HPC
 - ▶ Provides multiple communication APIs
 - ▶ Point-to-point
 - ▶ Partitioned point-to-point
 - ▶ RMA
 - ▶ Collective Communication (MPI_Allreduce, MPI_Bcast, etc.)

- ▶ HPC is used to solve large complex problems in many domains
 - ▶ Communication is one of the main bottlenecks in applications
- ▶ The Message Passing Interface (MPI)
 - ▶ Popular parallel programming model in HPC
 - ▶ Provides multiple communication APIs
 - ▶ Point-to-point
 - ▶ Partitioned point-to-point
 - ▶ RMA
 - ▶ Collective Communication (MPI_Allreduce, MPI_Bcast, etc.)
- ▶ MPI Threading Modes:
 - ▶ MPI_THREAD_SINGLE
 - ▶ MPI_THREAD_FUNNELLED
 - ▶ MPI_THREAD_SERIALIZED
 - ▶ MPI_THREAD_MULTIPLE

► MPI_THREAD_SERIALIZED

```
1 #pragma omp parallel for
2 for (int i=1; i<100; ++i)
3 {
4     /* Do Work */
5 }
6
7 MPI_Isend (...);
8 MPI_Irecv (...);
9
10 MPI_Waitall (...);
```


▶ MPI_THREAD_SERIALIZED

- ▶ All work is executed in the parallel region
- ▶ Wait for the all the threads to join
- ▶ Transfer Data

```
1 #pragma omp parallel for
2 for(int i=1; i<100; ++i)
3 {
4     /* Do Work */
5 }
6
7 MPI_Isend (...);
8 MPI_Irecv (...);
9
10 MPI_Waitall (...);
```

- ▶ MPI_THREAD_SERIALIZED
 - ▶ All work is executed in the parallel region
 - ▶ Wait for the all the threads to join
 - ▶ Transfer Data
- ▶ Some threads may finish early and may become idle as they wait

```
1 #pragma omp parallel for
2 for (int i=1; i<100; ++i)
3 {
4     /* Do Work */
5 }
6
7 MPI_Isend (...);
8 MPI_Irecv (...);
9
10 MPI_Waitall (...);
```

- ▶ `MPI_THREAD_SERIALIZED`
 - ▶ All work is executed in the parallel region
 - ▶ Wait for the all the threads to join
 - ▶ Transfer Data
- ▶ Some threads may finish early and may become idle as they wait
- ▶ Network bandwidth becomes under utilized.

```
1 #pragma omp parallel for
2 for(int i=1; i<100; ++i)
3 {
4     /* Do Work */
5 }
6
7 MPI_Isend (...);
8 MPI_Irecv (...);
9
10 MPI_Waitall (...);
```

► MPI_THREAD_MULTIPLE

```
1 #pragma omp parallel for
2 for (int i=1; i<100; ++i)
3 {
4     /* Do Work */
5     MPI_Isend (...);
6     MPI_Irecv (...);
7 }
8 MPI_Waitall (...);
```

► MPI_THREAD_MULTIPLE

- All work is executed in the parallel region
- As each thread completes its work, MPI calls are made

```
1 #pragma omp parallel for
2 for (int i=1; i<100; ++i)
3 {
4     /* Do Work */
5     MPI_Isend (...);
6     MPI_Irecv (...);
7 }
8 MPI_Waitall (...);
```

- ▶ MPI_THREAD_MULTIPLE
 - ▶ All work is executed in the parallel region
 - ▶ As each thread completes its work, MPI calls are made
- ▶ Although it may appear that data is transferred as soon as data is ready, communication may be serialized.

```
1 #pragma omp parallel for
2 for (int i=1; i<100; ++i)
3 {
4     /* Do Work */
5     MPI_Isend (...);
6     MPI_Irecv (...);
7 }
8 MPI_Waitall (...);
```

- ▶ `MPI_THREAD_MULTIPLE`
 - ▶ All work is executed in the parallel region
 - ▶ As each thread completes its work, MPI calls are made
- ▶ Although it may appear that data is transferred as soon as data is ready, communication may be serialized.
- ▶ Multi-threaded communication usually has issues with MPI's message matching queues.

```
1 #pragma omp parallel for
2 for (int i=1; i<100; ++i)
3 {
4     /* Do Work */
5     MPI_Isend (...);
6     MPI_Irecv (...);
7 }
8 MPI_Waitall (...);
```

MPI Partitioned Point-to-Point Communication

- ▶ MPI_Psend_init/MPI_Precv_init is used to initialize communication between processes.

```
1 MPI_Psend_init (...);
2 MPI_Precv_init (...);
3 MPI_Start (...);
4 #pragma omp parallel for
5 for (int i=1; i<100; ++i)
6 {
7     /* Do Work */
8     MPI_Pready (...);
9     int flag = 0;
10    while (!flag)
11    {
12        /* Do Work */
13        MPI_Parrived (...);
14    }
15    /* Do Work */
16 }
17 MPI_Waitall (...);
```


MPI Partitioned Point-to-Point Communication

- ▶ MPI_Psend_init/MPI_Precv_init is used to initialize communication between processes.
 - ▶ Message matching occurs here.
 - ▶ MPI Partitioned does not accept wildcards.

```
1 MPI_Psend_init (...);
2 MPI_Precv_init (...);
3 MPI_Start (...);
4 #pragma omp parallel for
5 for (int i=1; i<100; ++i)
6 {
7     /* Do Work */
8     MPI_Pready (...);
9     int flag = 0;
10    while (!flag)
11    {
12        /* Do Work */
13        MPI_Parrived (...);
14    }
15    /* Do Work */
16 }
17 MPI_Waitall (...);
```

MPI Partitioned Point-to-Point Communication

- ▶ MPI_Psend_init/MPI_Precv_init is used to initialize communication between processes.
 - ▶ Message matching occurs here.
 - ▶ MPI Partitioned does not accept wildcards.
- ▶ MPI_Start is called to start communication.

```
1 MPI_Psend_init (...);
2 MPI_Precv_init (...);
3 MPI_Start (...);
4 #pragma omp parallel for
5 for (int i=1; i<100; ++i)
6 {
7     /* Do Work */
8     MPI_Pready (...);
9     int flag = 0;
10    while (!flag)
11    {
12        /* Do Work */
13        MPI_Parrived (...);
14    }
15    /* Do Work */
16 }
17 MPI_Waitall (...);
```

MPI Partitioned Point-to-Point Communication

- ▶ MPI_Psend_init/MPI_Precv_init is used to initialize communication between processes.
 - ▶ Message matching occurs here.
 - ▶ MPI Partitioned does not accept wildcards.
- ▶ MPI_Start is called to start communication.
- ▶ A parallel for loop is launched

```
1 MPI_Psend_init (...);
2 MPI_Precv_init (...);
3 MPI_Start (...);
4 #pragma omp parallel for
5 for (int i=1; i<100; ++i)
6 {
7     /* Do Work */
8     MPI_Pready (...);
9     int flag = 0;
10    while (!flag)
11    {
12        /* Do Work */
13        MPI_Parrived (...);
14    }
15    /* Do Work */
16 }
17 MPI_Waitall (...);
```

MPI Partitioned Point-to-Point Communication

- ▶ MPI_Psend_init/MPI_Precv_init is used to initialize communication between processes.
 - ▶ Message matching occurs here.
 - ▶ MPI Partitioned does not accept wildcards.
- ▶ MPI_Start is called to start communication.
- ▶ A parallel for loop is launched
 - ▶ Work is Computed
 - ▶ Once data is ready, MPI_Pready is called.
 - ▶ Optionally, MPI_Parrived to check if incoming data has arrived.

```
1 MPI_Psend_init (...);
2 MPI_Precv_init (...);
3 MPI_Start (...);
4 #pragma omp parallel for
5 for (int i=1; i<100; ++i)
6 {
7     /* Do Work */
8     MPI_Pready (...);
9     int flag = 0;
10    while (!flag)
11    {
12        /* Do Work */
13        MPI_Parrived (...);
14    }
15    /* Do Work */
16 }
17 MPI_Waitall (...);
```

MPI Partitioned Point-to-Point Communication

- ▶ MPI_Psend_init/MPI_Precv_init is used to initialize communication between processes.
 - ▶ Message matching occurs here.
 - ▶ MPI Partitioned does not accept wildcards.
- ▶ MPI_Start is called to start communication.
- ▶ A parallel for loop is launched
 - ▶ Work is Computed
 - ▶ Once data is ready, MPI_Pready is called.
 - ▶ Optionally, MPI_Parrived to check if incoming data has arrived.
- ▶ MPI_Waitall is called to complete communication

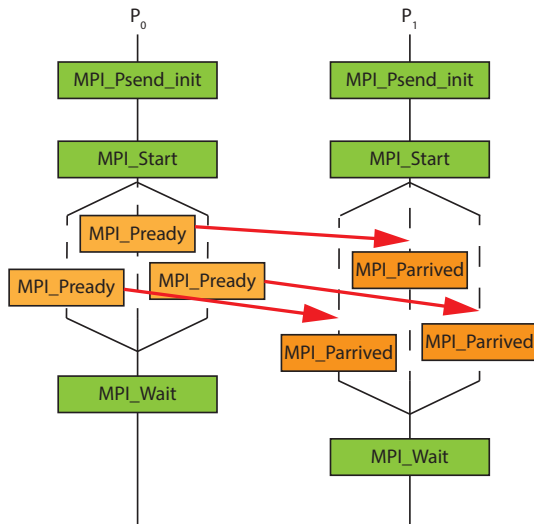
```
1 MPI_Psend_init (...);
2 MPI_Precv_init (...);
3 MPI_Start (...);
4 #pragma omp parallel for
5 for (int i=1; i<100; ++i)
6 {
7     /* Do Work */
8     MPI_Pready (...);
9     int flag = 0;
10    while (!flag)
11    {
12        /* Do Work */
13        MPI_Parrived (...);
14    }
15    /* Do Work */
16 }
17 MPI_Waitall (...);
```

MPI Partitioned Point-to-Point Communication

- ▶ MPI_Psend_init/MPI_Precv_init is used to initialize communication between processes.
 - ▶ Message matching occurs here.
 - ▶ MPI Partitioned does not accept wildcards.
- ▶ MPI_Start is called to start communication.
- ▶ A parallel for loop is launched
 - ▶ Work is Computed
 - ▶ Once data is ready, MPI_Pready is called.
 - ▶ Optionally, MPI_Parrived to check if incoming data has arrived.
- ▶ MPI_Waitall is called to complete communication
- ▶ A good implementation does not have the serialization issues of MPI Point-to-Point.

```
1 MPI_Psend_init (...);
2 MPI_Precv_init (...);
3 MPI_Start (...);
4 #pragma omp parallel for
5 for (int i=1; i<100; ++i)
6 {
7     /* Do Work */
8     MPI_Pready (...);
9     int flag = 0;
10    while (!flag)
11    {
12        /* Do Work */
13        MPI_Parrived (...);
14    }
15    /* Do Work */
16 }
17 MPI_Waitall (...);
```

MPI Partitioned Point-to-Point Communication



- ▶ Commonly used benchmarks do not support MPI Partitioned.
 - ▶ Sandia Micro Benchmarks (SMB)
 - ▶ OSU Micro Benchmarks (OSU)
 - ▶ Intel MPI Benchmarks (IMB)

- ▶ Commonly used benchmarks do not support MPI Partitioned.
 - ▶ Sandia Micro Benchmarks (SMB)
 - ▶ OSU Micro Benchmarks (OSU)
 - ▶ Intel MPI Benchmarks (IMB)
- ▶ Traditional point-to-point benchmarking techniques do not work for MPI Partitioned.

- ▶ Commonly used benchmarks do not support MPI Partitioned.
 - ▶ Sandia Micro Benchmarks (SMB)
 - ▶ OSU Micro Benchmarks (OSU)
 - ▶ Intel MPI Benchmarks (IMB)
- ▶ Traditional point-to-point benchmarking techniques do not work for MPI Partitioned.
- ▶ No production application uses MPI Partitioned
 - ▶ How can we discover possible candidates for porting?

- ▶ Commonly used benchmarks do not support MPI Partitioned.
 - ▶ Sandia Micro Benchmarks (SMB)
 - ▶ OSU Micro Benchmarks (OSU)
 - ▶ Intel MPI Benchmarks (IMB)
- ▶ Traditional point-to-point benchmarking techniques do not work for MPI Partitioned.
- ▶ No production application uses MPI Partitioned
 - ▶ How can we discover possible candidates for porting?

Research Questions

Can we design an MPI Partitioned Micro-benchmark to address the following:

- ▶ Commonly used benchmarks do not support MPI Partitioned.
 - ▶ Sandia Micro Benchmarks (SMB)
 - ▶ OSU Micro Benchmarks (OSU)
 - ▶ Intel MPI Benchmarks (IMB)
- ▶ Traditional point-to-point benchmarking techniques do not work for MPI Partitioned.
- ▶ No production application uses MPI Partitioned
 - ▶ How can we discover possible candidates for porting?

Research Questions

Can we design an MPI Partitioned Micro-benchmark to address the following:

- ▶ How can we understand the behaviour and performance of MPI Partitioned?

- ▶ Commonly used benchmarks do not support MPI Partitioned.
 - ▶ Sandia Micro Benchmarks (SMB)
 - ▶ OSU Micro Benchmarks (OSU)
 - ▶ Intel MPI Benchmarks (IMB)
- ▶ Traditional point-to-point benchmarking techniques do not work for MPI Partitioned.
- ▶ No production application uses MPI Partitioned
 - ▶ How can we discover possible candidates for porting?

Research Questions

Can we design an MPI Partitioned Micro-benchmark to address the following:

- ▶ How can we understand the behaviour and performance of MPI Partitioned?
- ▶ How could existing applications benefit from this new programming model?

- ▶ Commonly used benchmarks do not support MPI Partitioned.
 - ▶ Sandia Micro Benchmarks (SMB)
 - ▶ OSU Micro Benchmarks (OSU)
 - ▶ Intel MPI Benchmarks (IMB)
- ▶ Traditional point-to-point benchmarking techniques do not work for MPI Partitioned.
- ▶ No production application uses MPI Partitioned
 - ▶ How can we discover possible candidates for porting?

Research Questions

Can we design an MPI Partitioned Micro-benchmark to address the following:


- ▶ How can we understand the behaviour and performance of MPI Partitioned?
- ▶ How could existing applications benefit from this new programming model?
- ▶ What are appropriate partition sizes for application developers to use?

- ▶ Point-to-Point Metrics
 - ▶ Overhead
 - ▶ Perceived Bandwidth
 - ▶ Early Bird Communication
 - ▶ Availability

- ▶ Point-to-Point Metrics
 - ▶ Overhead
 - ▶ Perceived Bandwidth
 - ▶ Early Bird Communication
 - ▶ Availability
- ▶ Communication Patterns
 - ▶ Halo3D
 - ▶ Sweep3D

- ▶ Point-to-Point Metrics
 - ▶ Overhead
 - ▶ Perceived Bandwidth
 - ▶ Early Bird Communication
 - ▶ Availability
- ▶ Communication Patterns
 - ▶ Halo3D
 - ▶ Sweep3D
- ▶ Noise Models
 - ▶ Single Thread Delay
 - ▶ Uniform Distribution
 - ▶ Gaussian Distribution

- ▶ Point-to-Point Metrics
 - ▶ Overhead
 - ▶ Perceived Bandwidth
 - ▶ Early Bird Communication
 - ▶ Availability
- ▶ Communication Patterns
 - ▶ Halo3D
 - ▶ Sweep3D
- ▶ Noise Models
 - ▶ Single Thread Delay
 - ▶ Uniform Distribution
 - ▶ Gaussian Distribution
- ▶ Hot vs. Cold Cache



- ▶ Niagara Supercomputer at SciNet¹
 - ▶ 2x 20 Core Intel Skylake at 2.4GHz
 - ▶ EDR InfiniBand Network
 - ▶ GNU/Linux - CentOS 7.6
 - ▶ Open MPI (master branch)
 - ▶ UCX v1.11.0
 - ▶ MPIPCL

¹SciNet is funded by: the Canada Foundation for Innovation; the Government of Ontario; Ontario Research Fund - Research Excellence; and the University of Toronto. This research was enabled in part by support provided by the Digital Research Alliance of Canada

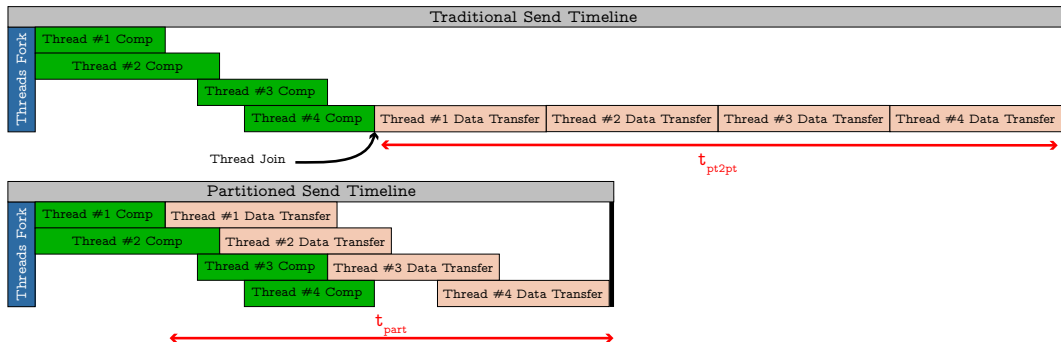
- ▶ What is the cost of using MPI Partitioned?

- ▶ What is the cost of using MPI Partitioned?
 - ▶ We measure each individual data transfer
 - ▶ Compare it to MPI Point-to-Point


$$Overhead = \frac{t_{part}}{t_{pt2pt}}$$

- ▶ What is the cost of using MPI Partitioned?
- ▶ We measure each individual data transfer
- ▶ Compare it to MPI Point-to-Point


$$\text{Overhead} = \frac{t_{\text{part}}}{t_{\text{pt2pt}}}$$



Overhead Results



(a) Cold Cache




(b) Hot Cache


Overhead of Partitioned Point-to-Point Communication Relative to Point-to-Point Communication for 10ms of Compute

Overhead Results

- Partition count correlates with overhead.



(a) Cold Cache




(b) Hot Cache


Overhead of Partitioned Point-to-Point Communication Relative to Point-to-Point Communication for 10ms of Compute

Overhead Results

- Partition count correlates with overhead.
- Overheads mostly impact small messages.



(a) Cold Cache



(b) Hot Cache

Overhead of Partitioned Point-to-Point Communication Relative to Point-to-Point Communication for 10ms of Compute

- ▶ What would be the required network bandwidth for MPI Point-to-Point to perform the same as MPI Partitioned?

Perceived Bandwidth



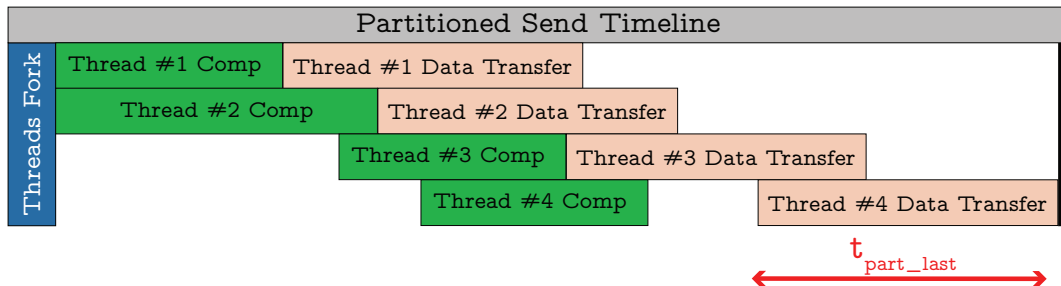
- What would be the required network bandwidth for MPI Point-to-Point to perform the same as MPI Partitioned?

$$\text{Perceived Bandwidth} = \frac{m}{t_{part_last}}$$


Perceived Bandwidth

- What would be the required network bandwidth for MPI Point-to-Point to perform the same as MPI Partitioned?


$$\text{Perceived Bandwidth} = \frac{m}{t_{\text{part_last}}}$$




Perceived Bandwidth Results




(a) 10ms Comp with 0% Noise




(b) 10ms Comp with 4% Noise




(c) 10ms Comp with 10% Noise



(d) 100ms Comp with 0% Noise



(e) 100ms Comp with 4% Noise




(f) 100ms Comp with 10% Noise


Perceived Bandwidth of MPI Partitioned Point-to-Point Communication with Uniform Noise and a Hot Cache for Different Noise and Compute Amounts

Perceived Bandwidth Results


- With 0% noise, we see our traditional bandwidth curve.




(a) 10ms Comp with 0% Noise




(b) 10ms Comp with 4% Noise




(c) 10ms Comp with 10% Noise



(d) 100ms Comp with 0% Noise



(e) 100ms Comp with 4% Noise




(f) 100ms Comp with 10% Noise


Perceived Bandwidth of MPI Partitioned Point-to-Point Communication with Uniform Noise and a Hot Cache for Different Noise and Compute Amounts

Perceived Bandwidth Results


- ▶ With 0% noise, we see our traditional bandwidth curve.
- ▶ Peak bandwidth is obtained for medium sized messages.




(a) 10ms Comp with 0% Noise




(b) 10ms Comp with 4% Noise




(c) 10ms Comp with 10% Noise



(d) 100ms Comp with 0% Noise



(e) 100ms Comp with 4% Noise




(f) 100ms Comp with 10% Noise


Perceived Bandwidth of MPI Partitioned Point-to-Point Communication with Uniform Noise and a Hot Cache for Different Noise and Compute Amounts

Perceived Bandwidth Results


- ▶ With 0% noise, we see our traditional bandwidth curve.
- ▶ Peak bandwidth is obtained for medium sized messages.
- ▶ Actual network bandwidth is saturated for large messages, thus perceived bandwidth drops.




(a) 10ms Comp with 0% Noise




(b) 10ms Comp with 4% Noise




(c) 10ms Comp with 10% Noise



(d) 100ms Comp with 0% Noise



(e) 100ms Comp with 4% Noise

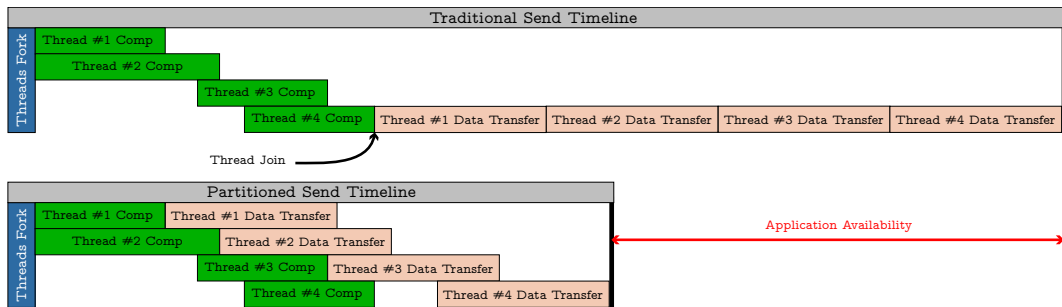



(f) 100ms Comp with 10% Noise

Perceived Bandwidth of MPI Partitioned Point-to-Point Communication with Uniform Noise and a Hot Cache for Different Noise and Compute Amounts


- ▶ If we switched to MPI Partitioned, how much extra work could the application do?

- If we switched to MPI Partitioned, how much extra work could the application do?





(a) Partition Counts (Single)




(b) Noise Distributions (16P)


Application Availability with 100ms Compute and 4% Noise

Application Availability

- Increasing partitions improve application availability for applications with a single thread delay.



(a) Partition Counts (Single)




(b) Noise Distributions (16P)


Application Availability with 100ms Compute and 4% Noise

Application Availability

- Increasing partitions improve application availability for applications with a single thread delay.
- Application noise will change how much can be gained from using MPI Partitioned.



(a) Partition Counts (Single)




(b) Noise Distributions (16P)


Application Availability with 100ms Compute and 4% Noise

Application Availability

- ▶ Increasing partitions improve application availability for applications with a single thread delay.
- ▶ Application noise will change how much can be gained from using MPI Partitioned.
- ▶ Less beneficial for very large messages.



(a) Partition Counts (Single)



(b) Noise Distributions (16P)

Application Availability with 100ms Compute and 4% Noise

Early Bird Communication



- ▶ How much communication occurs before the thread join?

Early Bird Communication



- ▶ How much communication occurs before the thread join?
 - ▶ How much do we overlap communication?

Early Bird Communication




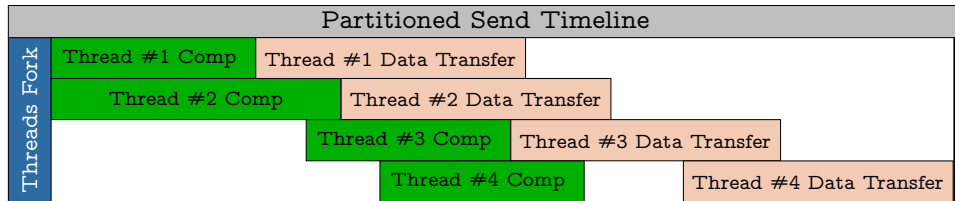
- ▶ How much communication occurs before the thread join?
 - ▶ How much do we overlap communication?

$$\% \text{ Early Bird} = \frac{t_{\text{before_join}}}{t_{\text{part}}}$$


Early Bird Communication

- How much communication occurs before the thread join?
 - How much do we overlap communication?


$$\% \text{ Early Bird} = \frac{t_{\text{before_join}}}{t_{\text{part}}}$$



Early Bird Communication




(a) 4% Noise




(b) 10% Noise

Early Bird Communication with 10ms Compute and Uniform Noise

Early Bird Communication



(a) 4% Noise




(b) 10% Noise


Early Bird Communication with 10ms Compute and Uniform Noise

Early Bird Communication

- Better overlap with more partitions for small message sizes.



(a) 4% Noise




(b) 10% Noise


Early Bird Communication with 10ms Compute and Uniform Noise

Early Bird Communication

- ▶ Better overlap with more partitions for small message sizes.
- ▶ More allows for more overlap for large messages.




(a) 4% Noise



(b) 10% Noise

Early Bird Communication with 10ms Compute and Uniform Noise


Sweep3D Communication Pattern



Sweep3D communication throughput for 16 partitions, 10ms compute, and 4% Single Noise with a Hot Cache

Sweep3D Communication Pattern


- Sweep3D communication pattern has lots of dependencies.



Sweep3D communication throughput for 16 partitions, 10ms compute, and 4% Single Noise with a Hot Cache

Sweep3D Communication Pattern


- Sweep3D communication pattern has lots of dependencies.
- Generally, multi-threading performs better than single threaded.



Sweep3D communication throughput for 16 partitions, 10ms compute, and 4% Single Noise with a Hot Cache

Sweep3D Communication Pattern


- Sweep3D communication pattern has lots of dependencies.
- Generally, multi-threading performs better than single threaded.
- The MPI Partition implementation in this paper is built upon MPI Send/Recv.



Sweep3D communication throughput for 16 partitions, 10ms compute, and 4% Single Noise with a Hot Cache

Sweep3D Communication Pattern


- ▶ Sweep3D communication pattern has lots of dependencies.
- ▶ Generally, multi-threading performs better than single threaded.
- ▶ The MPI Partition implementation in this paper is built upon MPI Send/Recv.
 - ▶ Therefore minimal difference for most message sizes.




Sweep3D communication throughput for 16 partitions, 10ms compute, and 4% Single Noise with a Hot Cache

Sweep3D Communication Pattern


- ▶ Sweep3D communication pattern has lots of dependencies.
- ▶ Generally, multi-threading performs better than single threaded.
- ▶ The MPI Partition implementation in this paper is built upon MPI Send/Recv.
 - ▶ Therefore minimal difference for most message sizes.
- ▶ Up to 15.1x higher throughput for large message sizes.



Sweep3D communication throughput for 16 partitions, 10ms compute, and 4% Single Noise with a Hot Cache



(a) 4 Partitions - 8 Threads




(b) 8 Partitions - 64 Threads (Oversubscribed)


Halo3D Communication Throughput For 10ms, 4% Single Noise with a Hot Cache

Halo3D Communication Pattern

- ▶ Halo exchange has lots of synchronization.
- ▶ Minimal difference in our different implementations.



(a) 4 Partitions - 8 Threads




(b) 8 Partitions - 64 Threads (Oversubscribed)


Halo3D Communication Throughput For 10ms, 4% Single Noise with a Hot Cache

Halo3D Communication Pattern

- ▶ Halo exchange has lots of synchronization.
 - ▶ Minimal difference in our different implementations.
- ▶ Additional work could be beneficial to Halo3D using MPI Partitioned
 - ▶ 64x increase in total computation with only a 16.8% decrease in throughput.
 - ▶ Could benefit from work stealing schemes.




(a) 4 Partitions - 8 Threads



(b) 8 Partitions - 64 Threads (Oversubscribed)


Halo3D Communication Throughput For 10ms, 4% Single Noise with a Hot Cache

Potential Application Improvements



Expected Speedup From Porting SNAP-C to
MPI Partitioned.


- The Sweep3D communication pattern showed potential for if it were ported to MPI Partitioned.



Expected Speedup From Porting SNAP-C to MPI Partitioned.

Potential Application Improvements

- ▶ The Sweep3D communication pattern showed potential for if it were ported to MPI Partitioned.
- ▶ SNAP uses a Sweep3D communication.
 - ▶ We profiled SNAP's communication.
 - ▶ Projected the potential speedup.



Expected Speedup From Porting SNAP-C to MPI Partitioned.

Conclusion And Future Work



- ▶ We provide the first MPI Partitioned Micro-Benchmark Suite.

Conclusion And Future Work



- ▶ We provide the first MPI Partitioned Micro-Benchmark Suite.
- ▶ Analyzed MPI Partitioned with a range of different metrics
 - ▶ Partition Count
 - ▶ Noise Distributions
 - ▶ Hot vs Cold Cache

Conclusion And Future Work



- ▶ We provide the first MPI Partitioned Micro-Benchmark Suite.
- ▶ Analyzed MPI Partitioned with a range of different metrics
 - ▶ Partition Count
 - ▶ Noise Distributions
 - ▶ Hot vs Cold Cache
- ▶ Showed what communication patterns could benefit from MPI Partitioned.

- ▶ We provide the first MPI Partitioned Micro-Benchmark Suite.
- ▶ Analyzed MPI Partitioned with a range of different metrics
 - ▶ Partition Count
 - ▶ Noise Distributions
 - ▶ Hot vs Cold Cache
- ▶ Showed what communication patterns could benefit from MPI Partitioned.

Future Work

- ▶ We provide the first MPI Partitioned Micro-Benchmark Suite.
- ▶ Analyzed MPI Partitioned with a range of different metrics
 - ▶ Partition Count
 - ▶ Noise Distributions
 - ▶ Hot vs Cold Cache
- ▶ Showed what communication patterns could benefit from MPI Partitioned.

Future Work

- ▶ Compare across different MPI implementations

- ▶ We provide the first MPI Partitioned Micro-Benchmark Suite.
- ▶ Analyzed MPI Partitioned with a range of different metrics
 - ▶ Partition Count
 - ▶ Noise Distributions
 - ▶ Hot vs Cold Cache
- ▶ Showed what communication patterns could benefit from MPI Partitioned.

Future Work

- ▶ Compare across different MPI implementations
- ▶ Porting Applications to MPI Partitioned

- ▶ We provide the first MPI Partitioned Micro-Benchmark Suite.
- ▶ Analyzed MPI Partitioned with a range of different metrics
 - ▶ Partition Count
 - ▶ Noise Distributions
 - ▶ Hot vs Cold Cache
- ▶ Showed what communication patterns could benefit from MPI Partitioned.

Future Work

- ▶ Compare across different MPI implementations
- ▶ Porting Applications to MPI Partitioned
- ▶ Accelerator-Triggered MPI Partitioned

- ▶ We provide the first MPI Partitioned Micro-Benchmark Suite.
- ▶ Analyzed MPI Partitioned with a range of different metrics
 - ▶ Partition Count
 - ▶ Noise Distributions
 - ▶ Hot vs Cold Cache
- ▶ Showed what communication patterns could benefit from MPI Partitioned.

Future Work

- ▶ Compare across different MPI implementations
- ▶ Porting Applications to MPI Partitioned
- ▶ Accelerator-Triggered MPI Partitioned
- ▶ MPI Partitioned Collectives

Acknowledgements



**Digital Research
Alliance** of Canada

**Alliance de recherche
numérique** du Canada